

Introduction

In this lab we will play with DDPM. We will choose our model, and our own noise schedule.

Implementation details

```
down_block_types=(
    "DownBlock2D", # a
    "DownBlock2D",
    "DownBlock2D",
    "DownBlock2D",
    "AttnDownBlock2D",
    "DownBlock2D",
),
up_block_types=(
    "UpBlock2D", # a
    "AttnUpBlock2D", #
    "UpBlock2D",
    "UpBlock2D",
    "UpBlock2D",
    "UpBlock2D",
),
```

I chose the hugging face diffuser UNet2DModel [1]. The down blocks I used are like the figure to the left. There are two layers per block and the channels are 128,128,256,256,512,512 from top to bottom.

The noise schedule is cosine_cap_v2 from hugging face's DDPM Scheduler and starts from beta 0.0001 to 0.02. The total steps T is set to 1000 for training.

I first resized the image to 64x64 to feed into the hugging face diffuser. Then, if I want to do evaluation, I have to first normalize the image to mean (0.5,0.5,0.5) and std (0.5,0.5,0.5) to feed into the evaluator.

For the class_label I used a linear layer to project the one-hot class encoding to the same dimension as the time embedding. Then I add the resulting vector with the time embedding.


The loss function is the same as the original DDPM paper.

$$L_{\text{simple}}(\theta) := \mathbb{E}_{t, \mathbf{x}_0, \epsilon} \left[\left\| \epsilon - \epsilon_{\theta}(\sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon, t) \right\|^2 \right]$$

Hyperparameters

I use a learning rate of 1e-5, AdamW as the optimizer. I trained for a total of 50 epochs.

Results and discussion



progressive generation



test.json

```
[pp037@ec037:~/DL/lab6]-[廿一時十八分50秒]-[G:class-embedding=]  
->$ python ddp.py --batch_size 20 --lr 0.00001 --save_img --ckpt checkpoints/28.pth --test (py11)  
loaded 18009 images from training  
100%|████████████████████████████████████████████████████████████████████████████████| 4/4 [02:56<00:00, 44.09s/it]  
100%|████████████████████████████████████████████████████████████████████████████████| 4/4 [05:51<00:00, 87.91s/it]  
average score 0.6864548494983278  
[pp037@ec037:~/DL/lab6]-[廿一時廿七分56秒]-[G:class-embedding=]  
->$ (py11)
```

test.json accuracy



new_test.json

```
[pp037@ec037:~/DL/Lab6]-[廿時廿七分五十六秒]-[G:class-embedding=]
->$ python ddp_m.py --batch-size 20 --lr 0.00001 --save_img --ckpt checkpoints/28.pth --test (py11)
loaded 18009 images from training
100%|██████████████████████████████████████████████████████████████████████████████| 4/4 [02:55<00:00, 43.78s/it]
0%|███████████████████████████████████████████████████████████████████████████████| | 0/4 [00:00<?, ?it/s]
100%|███████████████████████████████████████████████████████████████████████████████| 4/4 [05:50<00:00, 87.70s/it]
average score 0.7222886762360446
[pp037@ec037:~/DL/Lab6]-[廿時三十分二十秒]-[G:class-embedding=]
```

new_test.json accuracy

I originally fed the whole training image without resizing and got either a completely black image or a color-shifted image. Later someone told me I should resize the image first before feeding the image into the diffuser. I spent a lot of time wondering what I did wrong. I tried changing the model size, since the training time is unacceptable, changing to linear schedule, cosine schedule. Removing a couple of layers or reducing a couple of channels still produces a weird colored image. I even tried the DDIMScheduler but that is not a pure version of the original DDPM so I abandoned it.

Reference

- [1] <https://huggingface.co/docs/diffusers/api/models/unet2d>